

San Francisco Restaurant Inspection Data Analysis using SparkSQL and RDD's

General steps

- Create a case class for each data set
- Use CSV reader to read in each data file
- Convert RDD to DataFrame

Setting up input data sets

```
val baseDir = "/FileStore/tables/o7ra249x1502053895865/"
val raw_inspections = sc.textFile(s"$baseDir/inspections_plus.tsv")
val violations = sc.textFile(s"$baseDir/violations_plus.tsv")
val business = sc.textFile(s"$baseDir/businesses_plus.tsv")
```

```
baseDir: String = /FileStore/tables/o7ra249x1502053895865/
raw_inspections: org.apache.spark.rdd.RDD[String] = /FileStore/tables/o7ra249x1502053895865//inspections_plus.tsv MapPartitionsRDD[262] at textFile at <console>:36
violations: org.apache.spark.rdd.RDD[String] = /FileStore/tables/o7ra249x1502053895865//violations_plus.tsv MapPartitionsRDD[264] at textFile at <console>:37
business: org.apache.spark.rdd.RDD[String] = /FileStore/tables/o7ra249x1502053895865//businesses_plus.tsv MapPartitionsRDD[266] at textFile at <console>:38
```

1) What is the inspection score distribution like? (inspections_plus.csv)

Expected output - (*score*, *count*)

```
import org.apache.spark.sql.functions._
case class
Inspection(business_code:String,score:Int,date:String,comment:String)
val inspection=raw_inspections.map(_._split("\t"))
val validInspection=inspection.filter(l=>l(1) !=
"").map(p=>Inspection(p(0),p(1).trim.toInt,p(2),p(3)))
val InspectionDF=validInspection.toDF()
InspectionDF.groupBy("score").count().sort(desc("score")).show()
```

```
+-----+-----+
|score|count|
+-----+-----+
|  100| 3705|
|   98| 1534|
|   96| 2365|
|   94| 1751|
|   93|   374|
|   92| 1482|
|   91|   411|
|   90| 1241|
|   89|   480|
|   88|   640|
|   87|   440|
|   86|   483|
|   85|   415|
|   84|   319|
|   83|   328|
|   82|   250|
|   81|   289|
|   80|   266|
```

2) What is the risk category distribution like?

(violations_plus.csv)

Expected output - (*risk category, count*)

```
import org.apache.spark.sql.functions._
case class Risk(risk:String)
val risk=violations.map(_._split("\t")).filter(l=>l(3) != "N/A")
val validRisk=risk.map(l=>Risk(l(3)))
val RiskDF=validRisk.toDF()
RiskDF.groupBy("risk").count().show()
```

```

+-----+-----+
|          risk|count|
+-----+-----+
|    Low Risk|24717|
|Moderate Risk|15713|
|    High Risk| 6446|
+-----+-----+

```

```

import org.apache.spark.sql.functions._
defined class Risk
risk: org.apache.spark.rdd.RDD[Array[String]] = MapPartitionsRDD[179] at filter
at <console>:44
validRisk: org.apache.spark.rdd.RDD[Risk] = MapPartitionsRDD[180] at map at <co
nsole>:45
RiskDF: org.apache.spark.sql.DataFrame = [risk: string]

```

3) Which 20 businesses got lowest scores? (inspections_plus.csv, businesses_plus.csv)

(This should be more low score rather than lowest score)

Expected columns - (**business_id,name,address,city,postal_code,score**)

```

import org.apache.spark.sql.functions._
case class
Business(business_code:String,business_name:String,business_address:String,city
:String,postal_code:String)
val filtered_business=business.map(_._split("\t")).filter(l=>l(2)
!="").filter(l=>l(4).length ==5).map(p=>Business(p(0),p(1),p(2),p(3),p(4)))
val temp_businessDF=filtered_business.toDF()
//Replaces empty strings in the col city to San Francisco
val businessDF=temp_businessDF.na.replace("city", Map("-> "San Francisco"))
val
grouped_inspectionDF=InspectionDF.groupBy("business_code").agg(min("score").ali
as("score"))
//val joined=businessDF.join(grouped_inspectionDF, businessDF("business_code")
=== grouped_inspectionDF("business_code"), "inner") -- this will two cols of
business_code
val joined=businessDF.join(grouped_inspectionDF,Seq("business_code"))
joined.sort("score").limit(20).show()

```

```

+-----+-----+-----+-----+
----+-----+
|business_code|          business_name|          business_address|          city|postal_

```

code	score					
4133	42	Dick Lee Pastry	716 Jackson St	San Francisco		9
4133	46	ABC Bakery Cafe	650 Jackson St	San Francisco		9
4108	47	Imperial Palace	818 Washington St	San Francisco		9
4102	49	PUNJAB KABAB HOUSE	101 EDDY St	SF		9
4108	50	"Yummy Dim Sum & ..."	930 Stockton St	San Francisco		9
4111	51	Hong Kee & Kim	91 Drumm St	San Francisco		9
4133	51	New Asia Restaurant	772 Pacific Ave	San Francisco		9
		Yee's Restaurant	1131 Grant Ave	San Francisco		9

4) Which 20 businesses got highest scores?

(inspections_plus.csv, businesses_plus.csv)

Expected columns - (*business_id,name,address,city,postal_code,score*)

val

```
max_inspectionDF=InspectionDF.groupBy("business_code").agg(max("score").alias("score"))
```

```
val max_joined=businessDF.join(max_inspectionDF,"business_code")
```

```
max_joined.sort(desc("score")).limit(20).show()
```

business_code	business_name	business_address	city	postal_code	score
4122	Toyose INC	3814 Noriega St	Sf		9
4132	CAFE-ROSSO	1600 HOLLOWAY Ave	SF		9
4105	Chatz Coffee	215 02nd St	S.F.		9
4132	"Edith's Food Com..."	3251 20th Ave	San Francisco		9
	STARBUCK'S COFFEE...	2222 fillmore	SF		9

```

4115| 100|
|      5450|AT&T - (CART 3) ...|24 WILLIE MAYS PL...|      S.F.|      9
4107| 100|
|      36744|El Castillito Taq...| 250 Golden Gate Ave|San Francisco|      9
4102| 100|
|      24080|Cafe Mucol ...| 785 08th Ave ...|      S.F.|      9

```

5) Among all the restaurants that got 100 score, what kind of violations did they get (if any)

(inspections_plus.csv, violations_plus.csv)

(Examine "High Risk" violation only)

Expected columns - (**business_id**, **risk_category**, **date**, **description**)

Note - format the date in (**month/day/year**)

```

import org.apache.spark.sql.functions._
import org.apache.spark.sql.functions.{unix_timestamp, to_date,date_format}

case class
new_risk(business_code:String,date:String,violationTypeID:String,risk_category:
String,description:String)

val riskobj = violations.map(_._split("\t")).filter(l=>l(3) != "N/A").filter(l=>
(l(3)=="High Risk")).map(p=>new_risk(p(0),p(1),p(2),p(3),p(4)))
val new_riskDF=riskobj.toDF()

InspectionDF.createOrReplaceTempView("business")
val sqlDF = spark.sql("SELECT * FROM business WHERE score=100")

val joined_inspection_risk = new_riskDF.join(sqlDF,Seq("business_code","date"))
.select($"business_code",date_format(to_date(unix_timestamp($"date","yyyymmdd"))
.cast("timestamp"),"MM/dd/yyyy").alias("date"),$"risk_category",
  $"description").show()

```

```

+-----+-----+-----+-----+
|business_code|      date|risk_category|      description|
+-----+-----+-----+-----+
|      18825|01/21/2014|    High Risk|No hot water or r...|
|      1896|01/11/2014|    High Risk|Improper reheatin...|
|       17|01/23/2012|    High Risk|High risk food ho...|
|      5874|01/27/2014|    High Risk|Improper reheatin...|
|      3482|01/10/2015|    High Risk|Unclean or unsani...|

```

```
|          36744|01/26/2013|    High Risk|Improper cooling ...|
+-----+-----+-----+-----+
```

```
import org.apache.spark.sql.functions._
import org.apache.spark.sql.functions.{unix_timestamp, to_date, date_format}
defined class new_risk
riskobj: org.apache.spark.rdd.RDD[new_risk] = MapPartitionsRDD[222] at map at <
console>:65
new_riskDF: org.apache.spark.sql.DataFrame = [business_code: string, date: string ... 3 more fields]
sqlDF: org.apache.spark.sql.DataFrame = [business_code: string, score: int ... 2 more fields]
joined_inspection_risk: Unit = ()
```

6) Average inspection score by zip code

Expected columns - (**zip, average score with only two digits after decimal**)

```
//Find the avg for each restuarant
val
avg_inspectionDF=InspectionDF.groupBy("business_code").agg(avg("score").alias("score"))
val joined_business_inspection=
businessDF.join(avg_inspectionDF,"business_code").groupBy("postal_code").agg(round(avg($"score"),2).alias("score")).show(40)
```

```
+-----+-----+
|postal_code|score|
+-----+-----+
|          94102| 90.9|
|          94140| 96.5|
|          94107|94.76|
|          94104|93.57|
|          94131|93.21|
|          94014| 90.0|
|          94143| 89.5|
|          94609| 87.0|
|          94112|92.82|
|          94545|100.0|
|          92672|91.33|
|          94513|92.33|
|          94103|90.58|
|          94130|97.75|
|          94118|92.64|
|          94117| 91.6|
```

94129 83.0

7) Compute the proportion of all businesses in each neighborhood that have incurred at least one of the violations

- "High risk vermin infestation"
- "Moderate risk vermin infestation"
- "Sewage or wastewater contamination"
- "Improper food labeling or menu misrepresentation"
- "Contaminated or adulterated food"
- "Reservice of previously served foods"
- "Expected output: zip code, percentage"

This question is asking for each neighborhood, what is the proportion of businesses that have incurred at least one of the above nasty violations

Note: use UDF to determine which violations match with one of the above extremely bad violations

Expected columns - (***zip code, total violation count, extreme violation count, proportion with only two digits after decimal***)

```

import spark.implicits._
import org.apache.spark.sql.functions._
//create a udf
def extremeViolation(desc:String) : Int = {

    val violationList=List("High risk vermin infestation","Moderate risk vermin
infestation","Sewage or wastewater contamination","Improper food labeling or
menu misrepresentation","Contaminated or adulterated food","Reservice of
previously served foods")
    if (violationList.contains(desc)) return 1
    else return 0
}

//Initialize the udf
val ExtremeViolationUDF = udf(extremeViolation(_:String))

//Get the required fields from the tsv file
case class new_business(business_code:String,zip:String)
val new_value=business.map(_._split("\t")).filter(l=>l(4).length
==5).map(p=>new_business(p(0),p(4))).toDF()
case class new_violation(business_code:String,violations:String)
val
temp_value=violations.map(_._split("\t")).filter(l=>l.length==5).filter(l=>l(4)
!= "").map(p=>new_violation(p(0),p(4))).distinct().toDF()

val joined_table=new_value.join(temp_value,Seq("business_code"))
val res1=joined_table.groupBy($"zip").agg(count("*").alias("total violation
count"))
val
res2=joined_table.select($"zip",ExtremeViolationUDF($"violations").as("vc")).gr
oupBy($"zip").agg(round(sum($"vc"),2).alias("extreme violation count"))
val final_res=res1.join(res2,Seq("zip")).select($"zip","total violation
count","extreme violation count",round(((($"extreme violation count"/$"total
violation count")*100),2) as "proportion in percentage")

final_res.sort(desc("proportion in percentage")).show(100)

```

zip	total violation count	extreme violation count	proportion in percentage
94129	6	1	16.67
94115	1357	135	9.95
94109	2582	236	9.14
94143	22	2	9.09
94133	3291	299	9.09

94122	1948	171	8.78
94110	3142	255	8.12
94134	457	37	8.1
94116	602	47	7.81
94158	13	1	7.69
94114	1283	97	7.56
94108	1515	114	7.52
94117	1027	76	7.4
94102	2438	172	7.05
94112	1079	75	6.95
94123	1284	75	5.84
94124	730	41	5.62
94104	565	21	5.49

8) Are SF restaurants clean? Justify your answer

(Make to sure backup your answer with data - don't just state your opinion)

(Yes/No) and why

I feel SF restuarants are not extremely clean.If you see the violation for high and moderate risk, it is almost equal to 50% of the entire inspection dataset

